

Operating Systems & Systems Programming

Module 6

File System and Input Output Management

Dr. Vikash



Jaypee Institute of Information Technology, Noida



- 1 File Concepts
- 2 File Structure
- 3 Access Methods
- 4 Directory Structure
- 5 Protection
- 6 File-System Structure
- 7 Allocation Methods
- 8 Free Space Management



- To explain the function of file systems
- To describe the interfaces to file systems
- To discuss file-system design tradeoffs, including access methods, file sharing, file locking, and directory structures
- To explore file-system protection



- Contiguous logical address space
- Types:
 - Data (numeric, character, binary)
 - Program
- Contents defined by file's creator
 - Many types



- **Name**- only information kept in human-readable form
- **Identifier**- unique tag (number) identifies file within file system
- **Type**- needed for systems that support different types
- **Location**- pointer to file location on device
- **Size**- current file size
- **Protection**- controls who can do reading, writing, executing
- **Time, date, and user identification**- data for protection, security, and usage monitoring
- Information about files are kept in the directory structure, which is maintained on the disk
- Many variations, including extended file attributes such as file checksum
- Information kept in the directory structure



- File is an abstract data type
- **Create**
- **Write** - at write pointer location
- **Read** - at read pointer location
- **Reposition within file** - seek
- **Delete**
- **Truncate**
- **Open(Fi)**- search the directory structure on disk for entry Fi, and move the content of entry to memory
- **Close (Fi)**- move the content of entry Fi in memory to directory structure on disk


File Types - Name, Extension



file type	usual extension	function
executable	exe, com, bin or none	ready-to-run machine- language program
object	obj, o	compiled, machine language, not linked
source code	c, cc, java, pas, asm, a	source code in various languages
batch	bat, sh	commands to the command interpreter
text	txt, doc	textual data, documents
word processor	wp, tex, rtf, doc	various word-processor formats
library	lib, a, so, dll	libraries of routines for programmers
print or view	ps, pdf, jpg	ASCII or binary file in a format for printing or viewing
archive	arc, zip, tar	related files grouped into one file, sometimes com- pressed, for archiving or storage
multimedia	mpeg, mov, rm, mp3, avi	binary file containing audio or A/V information



- None - sequence of words, bytes
- Simple record structure
 - Lines
 - Fixed length
 - Variable length
- Complex Structures
 - Formatted document
 - Relocatable load file
- Can simulate last two with first method by inserting appropriate control characters
- Who decides:
 - Operating system
 - Program





- **Sequential Access**

- read next

- write next

- reset

- no read after last write (rewrite)

- **Direct Access-** file is fixed length **logical records**

- read n

- write n

- position to n

- read next

- write next

- rewrite n

n = relative block number

- Relative block numbers allow OS to decide where file should be placed

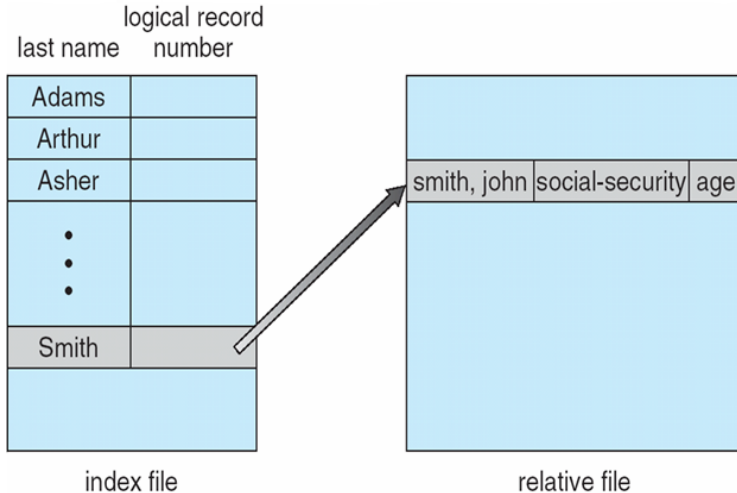


sequential access	implementation for direct access
<i>reset</i>	<i>cp = 0;</i>
<i>read next</i>	<i>read cp;</i> <i>cp = cp + 1;</i>
<i>write next</i>	<i>write cp;</i> <i>cp = cp + 1;</i>




- Can be built on top of base methods
- General involve creation of an index for the file
- Keep index in memory for fast determination of location of data to be operated on (consider UPC code plus record of data about that item)
- If too large, index (in memory) of the index (on disk)
- IBM indexed sequential-access method (ISAM)
 - Small master index, points to disk blocks of secondary index
 - File kept sorted on a defined key. All done by the OS
- VMS operating system provides index and relative files as another example

Example of Index and Relative Files






- A collection of nodes containing information about all files



- Both the directory structure and the files reside on disk

A Typical File-system Organization





- We mostly talk of general-purpose file systems
- But systems frequently have many file systems, some general- and some special- purpose
- Consider Solaris has
 - tmpfs - memory-based volatile FS for fast, temporary I/O
 - objfs - interface into kernel memory to get kernel symbols for debugging
 - ctfs - contract file system for managing daemons
 - lofs - loopback file system allows one FS to be accessed in place of another
 - procfs - kernel interface to process structures
 - ufs, zfs - general purpose file systems



- Search for a file
- Create a file
- Delete a file
- List a directory
- Rename a file
- Traverse the file system




The directory is organized logically to obtain

- Efficiency - locating a file quickly
- Naming - convenient to users
 - Two users can have same name for different files
 - The same file can have several different names

Grouping - logical grouping of files by properties, (e.g., all Java programs, all games, ...)




- A single directory for all users



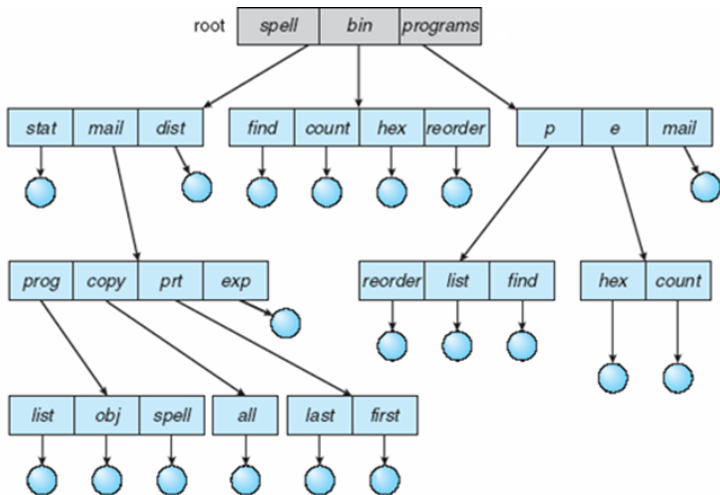
- Naming problem
- Grouping problem



- Separate directory for each user



- Path name
- Can have the same file name for different user
- Efficient searching
- No grouping capability

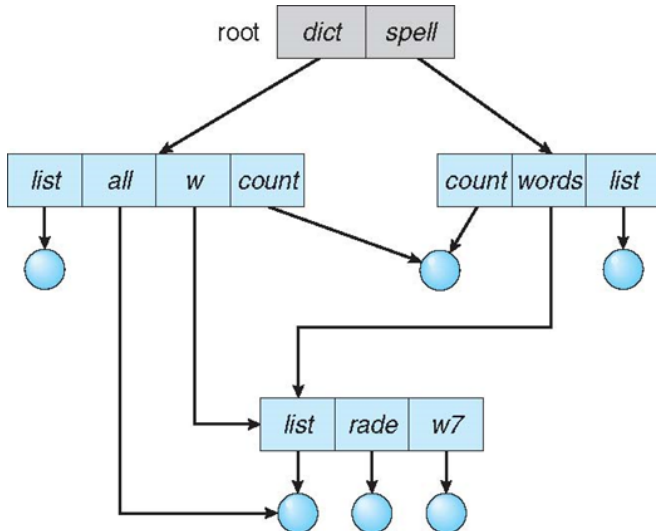




- Efficient Searching
- Grouping Capability
- Current directory (working directory)
 - **cd /spell/mail/prog**
 - **type list**



Have shared subdirectories and files

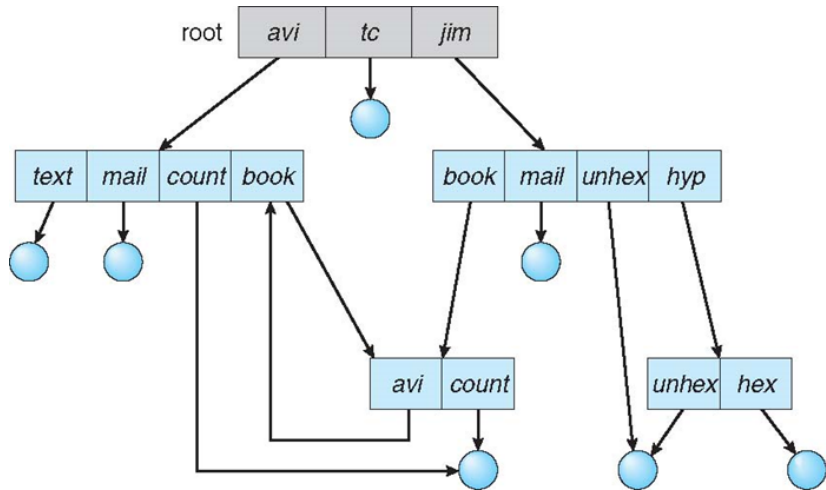




- Two different names (aliasing)
- If **dict** deletes **list** → dangling pointer

Solutions:

- Backpointers, so we can delete all pointers
- Backpointers using a daisy chain organization
- Entry-hold-count solution







- How do we guarantee no cycles?
 - Allow only links to file not subdirectories
 - **Garbage collection**
 - Every time a new link is added use a cycle detection algorithm to determine whether it is OK



- A file system must be **mounted** before it can be accessed
- A unmounted file system (i.e., Fig. 11-11(b)) is mounted at a **mount point**.



(a)



(b)



- Sharing of files on multi-user systems is desirable
- Sharing may be done through a **protection** scheme
- On distributed systems, files may be shared across a network
- Network File System (NFS) is a common distributed file-sharing method
- If multi-user system
 - **User IDs** identify users, allowing permissions and protections to be per-user
 - **Group IDs** allow users to be in groups, permitting group access rights
 - Owner of a file / directory
 - Group of a file / directory



- Uses networking to allow file system access between systems
 - Manually via programs like FTP
 - Automatically, seamlessly using **distributed file systems**
 - Semi automatically via the **world wide web**
- **Client-server** model allows clients to mount remote file systems from servers
 - Server can serve multiple clients
 - Client and user-on-client identification is insecure or complicated
 - **NFS** is standard UNIX client-server file sharing protocol
 - **CIFS** is standard Windows protocol
 - Standard operating system file calls are translated into remote calls
- Distributed Information Systems (**distributed naming services**) such as LDAP, DNS, NIS, Active Directory implement unified access to information needed for remote computing




- All file systems have failure modes
 - For example corruption of directory structures or other non-user data, called **metadata**
- Remote file systems add new failure modes, due to network failure, server failure
- Recovery from failure can involve **state information** about status of each remote request
- **Stateless** protocols such as NFS v3 include all information in each request, allowing easy recovery but less security



- File owner/creator should be able to control:
 - what can be done
 - by whom
- Types of access
 - **Read**
 - **Write**
 - **Execute**
 - **Append**
 - **Delete**
 - **List**




- Mode of access: read, write, execute
- Three classes of users on Unix / Linux
 - a. owner access 7 → 1 1 1
 - b. group access 6 → 1 1 0
 - c. public access 1 → 0 0 1
- Ask manager to create a group (unique name), say G, and add some users to the group.
- For a particular file (say game) or subdirectory, define an appropriate access.





- File structure
 - Logical storage unit
 - Collection of related information
- **File system** resides on secondary storage (disks)
 - Provided user interface to storage, mapping logical to physical
 - Provides efficient and convenient access to disk by allowing data to be stored, located retrieved easily
- Disk provides in-place rewrite and random access
 - I/O transfers performed in blocks of sectors (usually 512 bytes)
- **File control block** - storage structure consisting of information about a file
- **Device driver** controls the physical device
- File system organized into layers







- **Device drivers** manage I/O devices at the I/O control layer
 - Given commands like “read drive1, cylinder 72, track 2, sector 10, into memory location 1060” outputs low-level hardware specific commands to hardware controller
- **Basic file system** given command like retrieve block 123 translates to device driver
- Also manages memory buffers and caches (allocation, freeing, replacement)
 - Buffers hold data in transit
 - Caches hold frequently used data
- **File organization module** understands files, logical address, and physical blocks
- Translates logical block # to physical block #
- Manages free space, disk allocation



- **Logical file system** manages metadata information
 - Translates file name into file number, file handle, location by maintaining file control blocks (inodes in UNIX)
 - Directory management
 - Protection
- Layering useful for reducing complexity and redundancy, but adds overhead and can decrease performance
 - Translates file name into file number, file handle, location by maintaining file control blocks (inodes in UNIX)
 - Logical layers can be implemented by any coding method according to OS designer
- Many file systems, sometimes many within an operating system
 - Each with its own format (CD-ROM is ISO 9660; Unix has UFS, FFS; Windows has FAT, FAT32, NTFS as well as floppy, CD, DVD Blu-ray, Linux has more than 40 types, with extended file system ext2 and ext3 leading; plus distributed file systems, etc.)
 - New ones still arriving ZFS, GoogleFS, Oracle ASM, FUSE



- Mount table storing file system mounts, mount points, file system types
- The following figure illustrates the necessary file system structures provided by the operating systems
- Plus buffers hold data blocks from secondary storage
- Open returns a file handle for subsequent use *Data from read eventually copied to specified user process memory address*






- An allocation method refers to how disk blocks are allocated for files:
- **Contiguous allocation** - each file occupies set of contiguous blocks
 - Best performance in most cases
 - Simple – only starting location (block #) and length (number of blocks) are required
 - Problems include finding space for file, knowing file size, external fragmentation, need for
 - compaction off-line (downtime) or on-line



Mapping from logical to physical



Block to be accessed = $Q + \text{starting address}$

Displacement into block = R




- Many newer file systems (i.e., Veritas File System) use a modified contiguous allocation scheme
- Extent-based file systems allocate disk blocks in extents
- An extent is a contiguous block of disks
 - Extents are allocated for file allocation
 - A file consists of one or more extents



- **Linked allocation** – each file a linked list of blocks
 - File ends at nil pointer
 - No external fragmentation
 - Each block contains pointer to next block
 - No compaction, external fragmentation
 - Free space management system called when new block needed
 - Improve efficiency by clustering blocks into groups but increases internal fragmentation
 - Reliability can be a problem
 - Locating a block can take many I/Os and disk seeks
- **FAT (File Allocation Table) variation**
 - Beginning of volume has table, indexed by block number
 - Much like a linked list, but faster on disk and cacheable
 - New block allocation simple




Each file is a linked list of disk blocks: blocks may be scattered anywhere on the disk




Block to be accessed is the Qth block in the linked chain of blocks representing the file.

Displacement into block = $R + 1$




File-Allocation Table





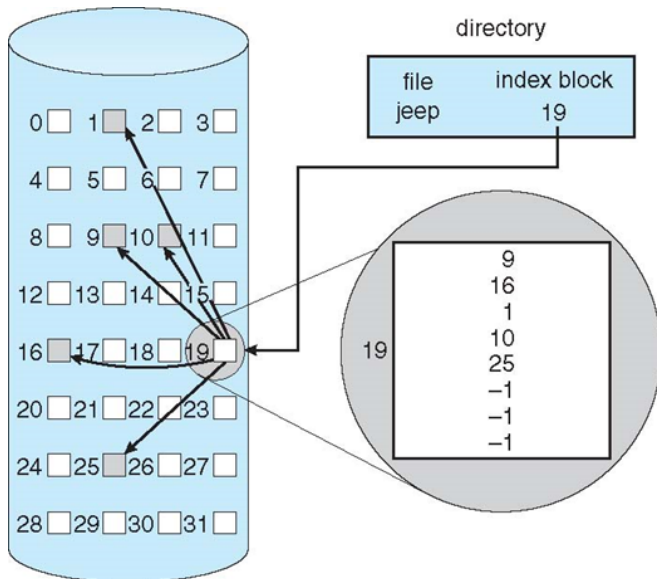
- **Indexed allocation**

- Each file has its own **index block(s)** of pointers to its data blocks
- Logical view




index table

Example of Indexed Allocation





- Need index table
- Random access
- Dynamic access without external fragmentation, but have overhead of index block
- Mapping from logical to physical in a file of maximum size of 256K bytes and block size of 512 bytes. We need only 1 block for index table




Q = displacement into index table

R = displacement into block




- Mapping from logical to physical in a file of unbounded length (block size of 512 words)
- Linked scheme Link blocks of index table (no limit on size)



Q_1 = block of index table

R_1 is used as follows:



Q_2 = displacement into block of index table

R_2 displacement into block of file:



- Two-level index (4K blocks could store 1,024 four-byte pointers in outer index - \therefore 1,048,567 data blocks and file size of up to 4GB)

$$LA / (512 \times 512) \begin{cases} Q_1 \\ R_1 \end{cases}$$


Q_1 = displacement into outer-index

R_1 is used as follows:

$$R_1 / 512 \begin{cases} Q_2 \\ R_2 \end{cases}$$

Q_2 = displacement into block of index table

R_2 displacement into block of file:





- File system maintains free-space list to track available blocks/clusters
 - (Using term “block” for simplicity)
- **Bit vector** or **bit map** (n blocks)




$$\text{bit}[i] = \begin{cases} 1 \Rightarrow \text{block}[i] \text{ free} \\ 0 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

Block Number Calculation

(number of bits per word) \times (number of 0-value words) + offset of first 1 bit



- Linked list (free list)
 - Cannot get contiguous space easily
 - No waste of space
 - No need to traverse the entire list (if # free blocks recorded)





- Grouping
 - Modify linked list to store address of next $n-1$ free blocks in first free block, plus a pointer to next block that contains free-block-pointers (like this one)
- Counting
 - Because space is frequently contiguously used and freed, with contiguous-allocation allocation, extents, or clustering
 - Keep address of first free block and count of following free blocks
 - Free space list then has entries containing addresses and counts



- Space Maps
 - Used in ZFS
 - Consider meta-data I/O on very large file systems
 - Full data structures like bit maps couldn't fit in memory → thousands of I/Os
 - Divides device space into metaslab units and manages metaslabs (Given volume can contain hundreds of metaslabs)
 - Each metaslab has associated space map (Uses counting algorithm)
 - But records to log file rather than file system (Log of all block activity, in time order, in counting format)
 - Metaslab activity → load space map into memory in balanced-tree structure, indexed by offset

Thank You!!!